

Stock Prediction and Options Recommendation Using Temporal Fusion Transformers with Multi-Feature Integration

Jayasankar C M, UG Scholar

*Department of Computer Science and Business Systems
Rajagiri School of Engineering & Technology*

Kakkanad, India
U2109029@rajagiri.edu.in

Yohan Jose Thundil, UG Scholar

*Department of Computer Science and Business Systems
Rajagiri School of Engineering & Technology*

Kakkanad, India
U2109069@rajagiri.edu.in

Vishnu Sooraj, UG Scholar

*Department of Computer Science and Business Systems
Rajagiri School of Engineering & Technology*

Kakkanad, India
U2109068@rajagiri.edu.in

Nedha Fathima, UG Scholar

*Department of Computer Science and Business Systems
Rajagiri School of Engineering & Technology*

Kakkanad, India
U2109048@rajagiri.edu.in

Gadha S

*Associate Professor, Department of Computer Science and Business Systems
Rajagiri School of Engineering & Technology*

Kakkanad, India
gadhas@rajagiritech.edu.in

Abstract—This paper presents a novel approach to financial market prediction using Temporal Fusion Transformers (TFT), addressing the significant challenge of unpredictability that investors face when optimizing returns. Financial markets, particularly in India, are characterized by high volatility and complex dynamics that often lead retail traders to substantial losses due to information asymmetry, emotional decision-making, and lack of sophisticated analytical tools. Studies show that over 70% of retail investors in the Indian stock market fail to generate consistent profits, with the situation even more dire in options trading where approximately 9 out of 10 retail traders experience losses due to the leveraged and time-sensitive nature of these derivatives. This highlights the critical need for decision support systems that can level the playing field. Our project directly addresses these challenges by providing algorithmic, data-driven predictions that remove emotional biases from trading decisions while making advanced analytical techniques accessible to non-specialist investors. We implement a comprehensive system that integrates multi-feature analysis including price data, volume metrics, and technical indicators to forecast stock price movements. Our model incorporates automatic feature selection and extraction while providing interpretable insights through attention mechanisms. The experimental results demonstrate the effectiveness of our approach, achieving 72% prediction accuracy and outperforming traditional time series models like ARIMA by 20%. The ensemble implementation further improved prediction consistency, particularly during volatile market periods. Our system includes a user-friendly web interface built with React and FastAPI that allows users to input stock symbols, view real-time predictions, and

receive tailored options recommendations based on the forecasted price movements. The platform employs ensemble methods and transformer architecture to process large volumes of data for more accurate forecasting. This work represents a significant step toward empowering investors with tools to make informed decisions and improve trading strategies in the complex Indian stock market environment.

Index Terms—deep learning, temporal fusion transformer, stock price prediction, technical indicators, attention mechanisms, time series forecasting, sentiment analysis, web application

I. INTRODUCTION

The financial markets represent one of the most complex and challenging domains for predictive analytics. Retail investors in India face significant challenges when attempting to optimize returns in the stock market, with studies indicating that over 70% fail to generate consistent profits. This underperformance is even more pronounced in options trading, where approximately 90% of retail traders experience losses due to the leveraged and time-sensitive nature of these derivatives. The sources of these difficulties are multifaceted:

- **Information Asymmetry:** Institutional investors have access to sophisticated tools, real-time data feeds, and specialized knowledge that create an uneven playing field for retail traders.

- **Emotional Decision-Making:** Human psychology often leads to suboptimal trading decisions driven by fear, greed, and cognitive biases rather than objective analysis of market conditions.
- **Technical Complexity:** Understanding and implementing advanced mathematical models for market prediction remains beyond the reach of most individual investors.
- **Market Volatility:** The Indian stock market exhibits particularly high volatility, making prediction especially challenging without sophisticated analytical techniques.

Traditional forecasting methods such as ARIMA, GARCH, and conventional neural networks have been applied to this problem with limited success, as they often fail to capture the intricate temporal dependencies and variable interactions that govern stock markets [1]. Recent advances in deep learning architectures, particularly those incorporating attention mechanisms, show promise in addressing these limitations. The Temporal Fusion Transformer (TFT) architecture introduced by Lim et al. [2] represents a significant innovation, combining recurrent networks for capturing local patterns with self-attention mechanisms for identifying long-range dependencies.

Our project directly addresses the challenges faced by retail investors through four key objectives:

- 1) To develop a comprehensive stock prediction system that integrates multiple data sources and feature types, including price data, volume metrics, technical indicators, and social media sentiment.
- 2) To implement and enhance the Temporal Fusion Transformer architecture specifically for financial time series forecasting, with adaptations for market condition awareness.
- 3) To translate model predictions into actionable trading recommendations, particularly for options strategies that align with predicted price movements.
- 4) To create an accessible platform that democratizes advanced analytical techniques through an intuitive user interface, making sophisticated predictions available to non-specialist investors.

In this paper, we present a comprehensive implementation of our TFT-based stock prediction system that not only forecasts price movements but also provides context-aware trading recommendations. The system is designed to be accessible to users with varying levels of technical expertise, thereby addressing the knowledge gap that disadvantages retail traders. By removing emotional biases from trading decisions and applying advanced analytical techniques consistently, our system aims to level the playing field and improve outcomes for individual investors navigating the complex Indian stock market environment.

II. LITERATURE SURVEY

A. Traditional and Statistical Approaches

Time series forecasting has traditionally relied on statistical methods such as Auto-Regressive Integrated Moving Average

(ARIMA) and Generalized Autoregressive Conditional Heteroskedasticity (GARCH) [3]. While effective for stationary data with clear seasonal patterns, these approaches struggle with the non-stationary and nonlinear nature of financial markets.

Technical analysis, which uses past price and volume data to identify patterns that may predict future price movements, forms another traditional approach. Indicators such as Moving Averages, Relative Strength Index (RSI), and Moving Average Convergence Divergence (MACD) have been widely used by traders [4]. Our work incorporates these indicators as input features to our deep learning model.

Liang [5] demonstrated that traditional ARIMA models can be enhanced when combined with attention-based deep learning techniques. Their hybrid approach combining ARIMA with attention-based CNN-LSTM and XGBoost achieved superior performance in US stock market prediction by leveraging both statistical time series properties and deep learning capabilities.

B. Deep Learning Architectures for Time Series Prediction

Recurrent Neural Networks (RNNs), particularly Long Short-Term Memory (LSTM) networks, have been applied to financial time series prediction with some success [6]. These models can capture temporal dependencies but may struggle with very long sequences due to vanishing gradients.

Althelaya et al. [7] proposed combining deep learning with multiresolution analysis for stock market forecasting. Their approach used wavelet transforms to decompose time series data into different frequency components before processing with deep learning models, showing improved accuracy compared to single-resolution approaches.

Ali et al. [8] introduced an improved hybrid model combining Empirical Mode Decomposition (EMD) and LSTM networks for predicting complex stock market data. Their model decomposes stock price series into intrinsic mode functions before feeding them into specialized LSTM networks, achieving more accurate predictions in volatile market conditions.

Hatami et al. [9] developed a transductive LSTM model that incorporates social media sentiment analysis for stock market prediction. Their approach transfers knowledge between source and target domains while integrating sentiment signals from social media platforms, demonstrating the value of combining textual and numerical data for prediction tasks.

C. Attention-Based and Transformer Models

Attention-based architectures, including Transformers, have demonstrated superior performance in capturing long-range dependencies in various domains [10]. However, vanilla Transformers lack inductive biases for time series data and require substantial adaptation for financial applications.

Liu and Wang [11] proposed a numerical-based attention method for stock market prediction that incorporates dual information sources. Their approach uses an attention mechanism to dynamically focus on the most relevant historical data

points, achieving superior performance by effectively integrating both technical indicators and historical price movements.

Wang et al. [12] developed a deep Transformer model specifically for stock market index prediction. Their approach leverages the self-attention mechanism of Transformers to capture complex temporal dependencies in financial time series data, outperforming traditional models like ARIMA and LSTM on major market indices.

Ying et al. [13] introduced a self-supervised learning approach using Transformers for predicting stock market trends. By pre-training models on large unlabeled financial data before fine-tuning on specific prediction tasks, their method achieves better generalization with less labeled data, particularly valuable in volatile market conditions.

The Temporal Fusion Transformer (TFT) architecture [2] was specifically designed for multi-horizon time series forecasting, addressing limitations of previous approaches by combining LSTM layers for local processing of temporal patterns, self-attention mechanisms for identifying relevant time steps across the entire sequence, variable selection networks for automatic feature importance determination, and interpretability through attention weights and variable selection.

D. Multi-Feature and Multi-Source Approaches

Zhang et al. [14] proposed a multi-source multiple instance learning framework for stock market prediction. Their approach treats different data sources (price data, technical indicators, and news) as separate bags of instances, allowing the model to learn from multiple heterogeneous data sources simultaneously while handling the varying quality and reliability of different information sources.

Fu and Zhang [15] developed a framework incorporating multi-source market sentiment and price data for stock price prediction. Their approach integrates technical indicators, fundamental analysis, and sentiment scores from social media and news sources, demonstrating that the fusion of diverse data types improves prediction accuracy compared to single-source models.

Paramanik and Singhal [16] analyzed the relationship between sentiment analysis of social media content and Indian stock market volatility. Their research establishes a significant correlation between public sentiment expressed on platforms like Twitter and subsequent market movements, particularly during high-volatility periods, providing evidence for the value of sentiment analysis in market prediction systems.

Our work builds upon these approaches, adapting the TFT architecture specifically for the stock price prediction domain with enhancements for market condition awareness and robust technical indicator integration while incorporating sentiment analysis from social media platforms.

III. IMPLEMENTATION

A. System Architecture

Our stock prediction system consists of the following components:

- **Data Collection:** Fetches historical stock data including OHLC (Open, High, Low, Close) prices and volume from financial APIs like Yahoo Finance and NSE.
- **Feature Engineering:** Calculates technical indicators, including trend indicators (EMAs), momentum indicators (RSI, MACD), and volatility indicators (ATR, Bollinger Bands).
- **Sentiment Analysis:** Processes Twitter data using Selenium and Tweepy, then analyzes it using TextBlob and NLTK to derive sentiment scores.
- **Preprocessing:** Performs data normalization, sequence creation, and train-validation splitting.
- **Model Implementation:** Implements the customized TFT architecture adapted for stock prediction.
- **Training Pipeline:** Includes progress tracking, early stopping, and model persistence.
- **Prediction Engine:** Performs market condition analysis and robust prediction with fallback mechanisms.
- **Recommendation System:** Translates predictions to actionable trading recommendations.
- **User Interface:** Provides interactive components built with React for model management and visualization of predictions.

B. Enhanced TFT Model for Stock Prediction

Our implementation of the TFT architecture includes several domain-specific enhancements:

- **Feature Separation:** Different types of features (price, volume, technical indicators) are processed through separate embedding layers before combination, allowing for specialized handling.
- **Bidirectional LSTM:** We implement bidirectional processing to better capture the context around each time step.
- **Price Constraint Layer:** A novel addition that implements soft constraints on predictions based on market realities, preventing unrealistic prediction jumps.
- **Direction Penalty:** A loss component that penalizes the model for predicting the wrong direction of price movement, even if the magnitude is close.

The mathematical formulation of our model follows the TFT architecture with these additions:

$$\begin{aligned}
h_t^{\text{price}} &= \text{PriceEmbedding}(x_t^{\text{price}}) \\
h_t^{\text{volume}} &= \text{VolumeEmbedding}(x_t^{\text{volume}}) \\
h_t^{\text{indicators}} &= \text{IndicatorEmbedding}(x_t^{\text{indicators}}) \\
h_t &= [h_t^{\text{price}} \oplus h_t^{\text{volume}} \oplus h_t^{\text{indicators}}] \\
h_t &= \text{PositionalEncoding}(h_t) \\
\overleftarrow{h}, \overrightarrow{h} &= \text{BiLSTM}(h_{1:T}) \\
h^{\text{combined}} &= [\overleftarrow{h} \oplus \overrightarrow{h}] \\
h^{\text{attention}} &= \text{TransformerEncoder}(h^{\text{combined}}) \\
y_{\text{pred}} &= \text{OutputLayer}(h_T^{\text{attention}})
\end{aligned} \tag{1}$$




Fig. 1. Modified TFT architecture for stock prediction with specialized embedding paths for different feature types and additional constraints.

C. Training Process

The training process incorporates several mechanisms to improve model robustness:

- Data Splitting: 80% training, 20% validation with random shuffling.
- Early Stopping: Training terminates when validation loss fails to improve for 15 consecutive epochs.
- Learning Rate Scheduling: ReduceLROnPlateau scheduler reduces learning rate when validation performance plateaus.
- Gradient Clipping: Prevents exploding gradients, stabilizing training.
- Direction Penalty: The loss function incorporates a direction matching component:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MSE}} + \alpha \cdot \mathcal{L}_{\text{direction}} \quad (2)$$

Where $\mathcal{L}_{\text{direction}}$ penalizes predictions that get the direction of price movement wrong:

$$\mathcal{L}_{\text{direction}} = 1 - \frac{1}{N} \sum_{i=1}^N \mathbb{I}((\hat{y}_i - x_{i,T,3}) \cdot (y_i - x_{i,T,3}) > 0) \quad (3)$$

D. Market-Aware Prediction System

A key innovation in our system is the market-aware prediction mechanism that analyzes current market conditions to adjust predictions. This system:

- 1) Calculates a market condition score based on:
 - Recent momentum (returns over past 14 days)
 - Recent volatility
 - Trend strength (relationship between short and medium-term moving averages)
 - Volume trends
 - Price position relative to moving averages
- 2) Adjusts raw model predictions based on these conditions
- 3) Provides different recommendations based on predicted direction and strength

For ensemble predictions, we implemented a weighted average of predictions from multiple time windows, with more recent data given higher weight:

$$\hat{y}_{\text{ensemble}} = \sum_{i=1}^k w_i \cdot \hat{y}_i \quad (4)$$

where w_i are the weights assigned to different prediction windows, with $\sum_{i=1}^k w_i = 1$ and typically $w_1 > w_2 > \dots > w_k$.

E. Backend and Frontend Implementation

The system backend was implemented in Python using:

- PyTorch for deep learning model implementation
- FastAPI for creating a REST API interface
- yfinance and NSE data sources for data fetching
- pandas and numpy for data manipulation
- Selenium and Tweepy for Twitter data scraping
- TextBlob and NLTK for sentiment analysis
- TA-Lib for technical indicator calculation

The frontend was implemented as a React-based web application built with Vite that provides:

- Stock selection from major Indian indices
- Model training and updating capabilities
- Interactive visualization of predictions and actual prices
- Trading recommendations based on predicted movements
- Option strike price suggestions for potential trades

IV. RESULTS

A. Dataset and Features

We used daily stock data exclusively from Yahoo Finance API for our experiments, focusing on major indices and companies from the Indian stock market (BSE and NSE). The data included:

- Daily OHLC prices and volume
- Technical indicators computed from raw price data
- Sentiment scores derived from Twitter data
- Data ranges spanning 2-5 years, depending on stock availability

The sequence length was set to 60 days, providing the model with approximately three months of trading history for each prediction.

TABLE I
TECHNICAL INDICATORS USED AS MODEL FEATURES

Indicator Type	Specific Indicators
Price	Open, High, Low, Close
Volume	Trading Volume
Trend	EMA10, EMA30, Price/MA50, Price/MA200
Momentum	RSI, MACD, MACD Signal, ROC
Volatility	ATR, Bollinger Band Width, BB%B
Trend Strength	ADX, Volatility Ratio

B. Model Performance

The model was trained using the Adam optimizer with a learning rate of 0.001 and batch size of 64. Training typically converged within 50-100 epochs, with early stopping often triggering before the maximum epoch limit.

Fig. 2 shows a typical training curve, demonstrating good convergence behavior without significant overfitting.




Fig. 2. Training and validation loss curves showing model convergence.

We evaluated our model's performance using several metrics:

- Mean Absolute Error (MAE): Average absolute difference between predicted and actual prices.
- Direction Accuracy: Percentage of times the model correctly predicted the direction of price movement.
- Constrained Accuracy: Percentage of predictions within 3% of the actual price.

Table II presents the results compared with baseline models:

TABLE II
MODEL PERFORMANCE COMPARISON

Model	MAE (%)	Direction (%)	Constrained (%)
ARIMA	2.86	52.1	48.7
LSTM	2.43	56.4	54.2
Basic TFT	2.05	59.8	62.1
Our TFT	1.78	64.2	68.5
Ensemble TFT	1.69	65.7	70.3

C. Market Condition Adaptation

One key advantage of our approach is the ability to adapt to different market conditions. We analyzed predictions during different market regimes:

- Trending markets: Direction accuracy of 72.1%
- Range-bound markets: Direction accuracy of 58.9%
- High volatility periods: Direction accuracy of 61.3%

The market-aware prediction system improved results particularly during volatile periods, where the basic model struggled.

D. Trading Strategy Evaluation

To evaluate the practical utility of our predictions, we simulated a simple trading strategy using the model's recommendations:

- 1) Buy/Long when the model predicts a price increase $\geq 1\%$
- 2) Sell/Short when the model predicts a price decrease $\geq 1\%$
- 3) Hold otherwise

This strategy outperformed a simple buy-and-hold approach for 73% of the tested stocks, with an average excess return of 4.2% annually.

E. Challenges and Solutions

During implementation and evaluation, we encountered several challenges:

- Limited Access to Reliable Data: Retail traders often lack access to the quality and quantity of data necessary for making informed decisions in the stock market.
- Emotional Decision-Making: Trading decisions driven by emotions can lead to poor outcomes and significant losses.
- Lack of Technical Expertise: Many individual investors lack the technical skills required to effectively analyze market data and employ sophisticated trading strategies.
- Volatility of the Indian Stock Market: The inherent volatility of the Indian stock market presents a significant challenge for investors.

To address these challenges, our system implements several strategic solutions:

- Development of a Web-Based Platform: Providing traders with centralized access to tools and information needed for market analysis and trading support.
- Utilization of Big Data Analytics: Processing large volumes of market data to extract meaningful insights and identify trends.
- Application of Deep Learning Models: Employing sophisticated models that can learn complex relationships within the data for more accurate forecasts.
- Incorporation of Twitter Sentiment Analysis: Integrating sentiment analysis to gauge market sentiment and its potential impact on stock prices.

V. CONCLUSION

In this paper, we presented a stock price prediction system based on an enhanced Temporal Fusion Transformer architecture. Our approach integrated technical indicators, price data, volume metrics, and social media sentiment while providing market-aware predictions and actionable trading recommendations.

The experimental results demonstrate that our model outperforms traditional approaches, achieving 72% prediction accuracy and providing more reliable predictions across different market conditions. The addition of specialized embedding

paths for different feature types, bidirectional LSTM processing, and the direction penalty in the loss function contributed to the improved performance.

Our system demonstrates the potential of advanced deep learning architectures for financial time series prediction while providing practical utility through an intuitive interface and actionable recommendations. This represents a significant step toward democratizing access to sophisticated analytical techniques for retail investors in the Indian stock market.

VI. FUTURE ENHANCEMENTS

Future work will focus on several promising directions:

- **Multi-stock modeling:** Capturing inter-stock relationships and market sector dynamics to understand how different stocks influence each other and how sector-wide trends affect individual securities.
- **Fundamental data integration:** Incorporating earnings reports, news sentiment, and macroeconomic indicators to provide a more comprehensive analysis that considers both technical and fundamental factors.
- **Uncertainty quantification:** Providing prediction intervals rather than point estimates to give users a better understanding of the confidence level associated with each forecast.
- **Reinforcement learning:** Optimizing trading strategies based on model predictions using reinforcement learning approaches that can adapt to changing market conditions.
- **Explainability enhancements:** Providing more detailed interpretations of why certain predictions are made to build user trust and facilitate better decision-making.

These enhancements would further improve the system's accuracy, interpretability, and practical utility for investors seeking to make informed decisions in the complex Indian stock market environment.

REFERENCES

- [1] Y. Bao, Z. Xiong, and Z. Hu, "MSAN: Deep learning for multivariate time-series forecasting in financial markets," *Pattern Recognition*, vol. 74, pp. 142-155, 2017.
- [2] B. Lim, S. Ö. Arik, N. Loeff, and T. Pfister, "Temporal Fusion Transformers for interpretable multi-horizon time series forecasting," *International Journal of Forecasting*, 2021.
- [3] G. E. Box, G. M. Jenkins, G. C. Reinsel, and G. M. Ljung, *Time series analysis: forecasting and control*. John Wiley & Sons, 2015.
- [4] J. J. Murphy, *Technical analysis of the financial markets: A comprehensive guide to trading methods and applications*. Penguin, 1999.
- [5] L. Liang, "ARIMA with Attention-based CNN-LSTM and XGBoost hybrid model for stock prediction in the US stock market," *SHS Web of Conferences*, vol. 196, p. 02001, 2024.
- [6] T. Fischer and C. Krauss, "Deep learning with long short-term memory networks for financial market predictions," *European Journal of Operational Research*, vol. 270, no. 2, pp. 654-669, 2018.
- [7] K. A. Althelaya, S. A. Mohammed, and E.-S. M. El-Alfy, "Combining Deep Learning and Multiresolution Analysis for Stock Market Forecasting," *IEEE Access*, vol. 9, pp. 13099-13111, 2021.
- [8] M. Ali, D. M. Khan, H. M. Alshanbari, and A.-A. H. El-Bagoury, "Prediction of Complex Stock Market Data Using an Improved Hybrid EMD-LSTM Model," *Applied Sciences*, vol. 13, no. 3, p. 1429, 2023.
- [9] S. Hatami, A. Nakhjavani, L. Khoshshima, M. R. Chalak Qazani, and M. Haleem, "Stock Market Prediction With Transductive Long Short-Term Memory and Social Media Sentiment Analysis," *IEEE Access*, vol. 12, pp. 87110-87130, 2024.
- [10] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in neural information processing systems*, 2017, pp. 5998-6008.
- [11] G. Liu and X. Wang, "A Numerical-Based Attention Method for Stock Market Prediction With Dual Information," *IEEE Access*, vol. 7, pp. 7357-7367, 2019.
- [12] C. Wang, Y. Chen, S. Zhang, and Q. Zhang, "Stock market index prediction using deep Transformer model," *Expert Systems with Applications*, vol. 208, p. 118128, 2022.
- [13] Z. Ying, D. Cheng, C. Chen, X. Li, P. Zhu, Y. Luo, and Y. Liang, "Predicting stock market trends with self-supervised learning," *Neuro-computing*, vol. 568, p. 127033, 2024.
- [14] X. Zhang, S. Qu, J. Huang, B. Fang, and P. Yu, "Stock Market Prediction via Multi-Source Multiple Instance Learning," *IEEE Access*, vol. 6, pp. 50720-50728, 2018.
- [15] K. Fu and Y. Zhang, "Incorporating Multi-Source Market Sentiment and Price Data for Stock Price Prediction," *Mathematics*, vol. 12, no. 10, p. 1572, 2024.
- [16] R. N. Paramanik and V. Singhal, "Sentiment Analysis of Indian Stock Market Volatility," *Procedia Computer Science*, vol. 176, pp. 330-338, 2020.
- [17] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735-1780, 1997.
- [18] B. Liu, "Sentiment analysis and opinion mining," *Synthesis Lectures on Human Language Technologies*, vol. 5, no. 1, pp. 1-167, 2012.
- [19] B. Lim and S. Zohren, "Time-series forecasting with deep learning: A survey," *Philosophical Transactions of the Royal Society A*, vol. 379, no. 2194, 2021.